

# Probability machines

## Risk estimation in a post-parameter world

Abhijit Dasgupta and James Malley



# Effect size

Let's go back to your first regression course

*"How much does the outcome change, on average, when a predictor changes by one unit, all other predictors remaining the same?"*

# Effect size

Let's go back to your first regression course

*"How much does the outcome change, on average, when a predictor changes by one unit, all other predictors remaining the same?"*

This is based on the concept of counterfactuals

- What would happen **if** a predictor changed by 1 unit keeping everything else the same?
- This is not something we can observe, but only something we can conceptualize
- Using multiple machines, we can actually estimate these counterfactuals directly **for each observation!!!**
- We don't need parameters in the model; we can do this nonparametrically
- Requires us to think about what exactly are the effects we want to interrogate

# Counterfactuals

X = 1



X = 0

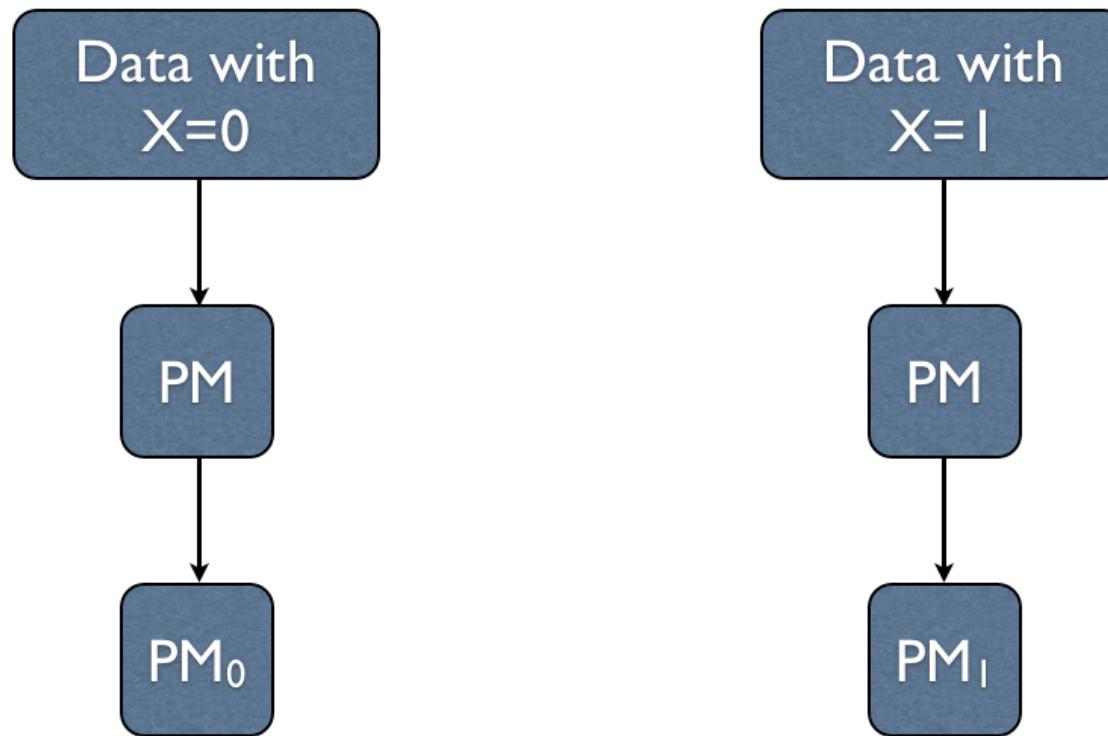


The counterfactual argument is, in essence

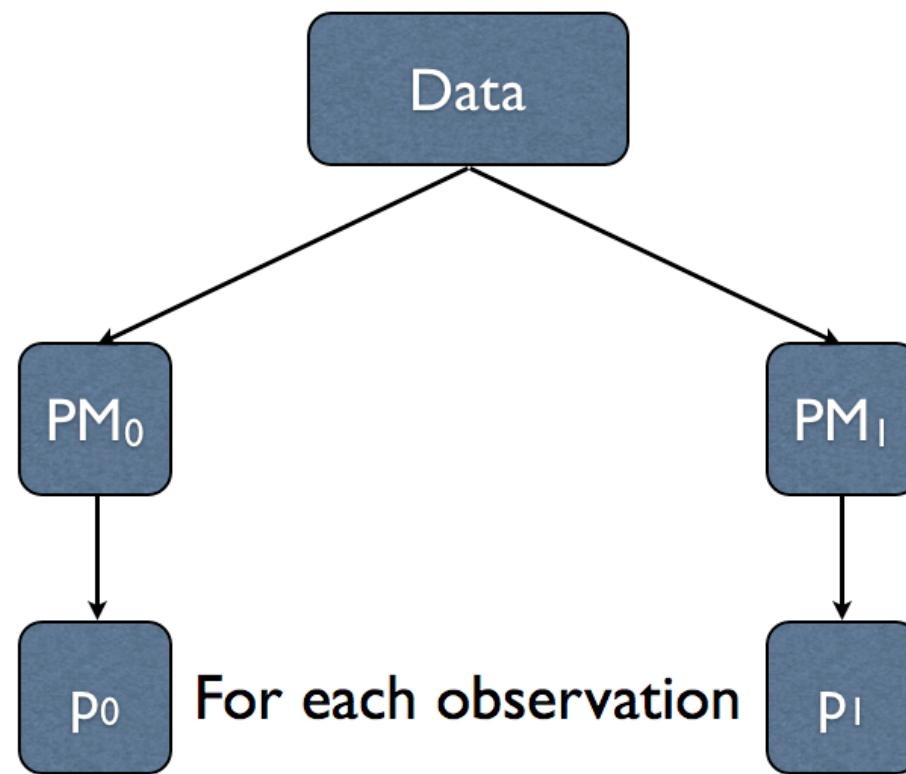
If we put an observation in the other landscape, what would it do?

# Counterfactual machine

# Counterfactual machine



# Counterfactual machine



# Risk machine

- $PM_1$  captures the landscape for  $X=1$
- $PM_0$  captures the landscape for  $X=0$

Now put each observation in each landscape and record its predicted outcome

Note, for each observation we now have a  $p_1$  and a  $p_0$

Now we can compute *conditional odds ratios* using

$$OR = \frac{p_1(1 - p_0)}{(1 - p_1)p_0}$$

for each observation, and look at group-specific odds ratios by averaging or taking medians

# Simulations

We generate data from a logistic regression model with

- 10 independent binary features
- 3 features associated with outcome to various degrees
- 7 features not associated with outcome (to mimic sparseness)

# Simulations

We fit three models to the generated data

- Main effects logistic regression

```
glm(y~x1+x2+x3+..., family=binomial)
```

- Main effects + two-way interactions logistic regression

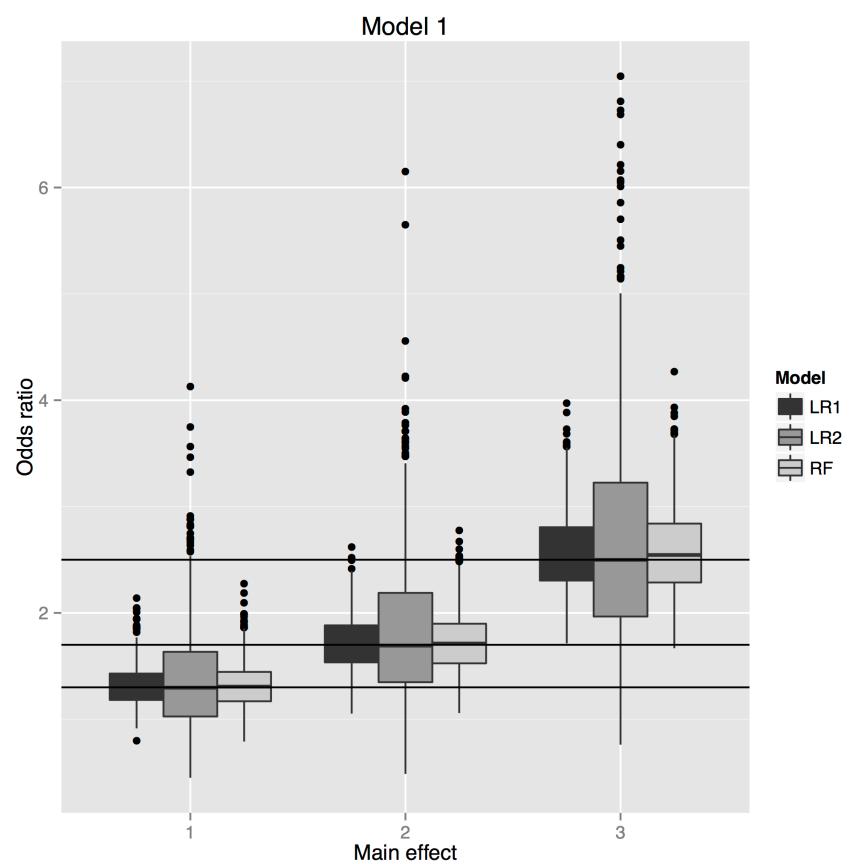
```
glm(y~(x1+x2+x3+...)^2, family=binomial)
```

- Random forest regression

```
randomForest(y~x1+x2+x3+...)
```

For this entire exercise, **we do not change this code**

# Simulations



Look at a main effects model

- All 3 are unbiased
- RF does as well as logistic regression for efficiency

# WARNING: Soapbox time

# Logistic regression and odds ratios

Logistic regression (the industry standard) gives us odds ratios (the industry standard)

# Logistic regression and odds ratios

Logistic regression (the industry standard) gives us odds ratios (the industry standard)

Please try explaining what an odds ratio is

# Logistic regression and odds ratios

Logistic regression (the industry standard) gives us odds ratios (the industry standard)

- The odds ratio is interpreted as a risk ratio
  - Can only do this when the outcome is rare
- The odds ratio might be reasonable for gamblers (Hello, Bernoullis), but probably not for clinicians
- If I had my way, we'd report either risk ratios or risk differences

Now back to regular programming

# Risk machine

We have individual  $p_1$  and  $p_0$ , so we can directly compute

- risk differences

$$RD = p_1 - p_0$$

- risk ratios

$$RR = p_1/p_0$$

# Risk machine

## What about interactions?

Since we have a way of estimating counterfactuals, estimating conditional interaction effects are straightforward

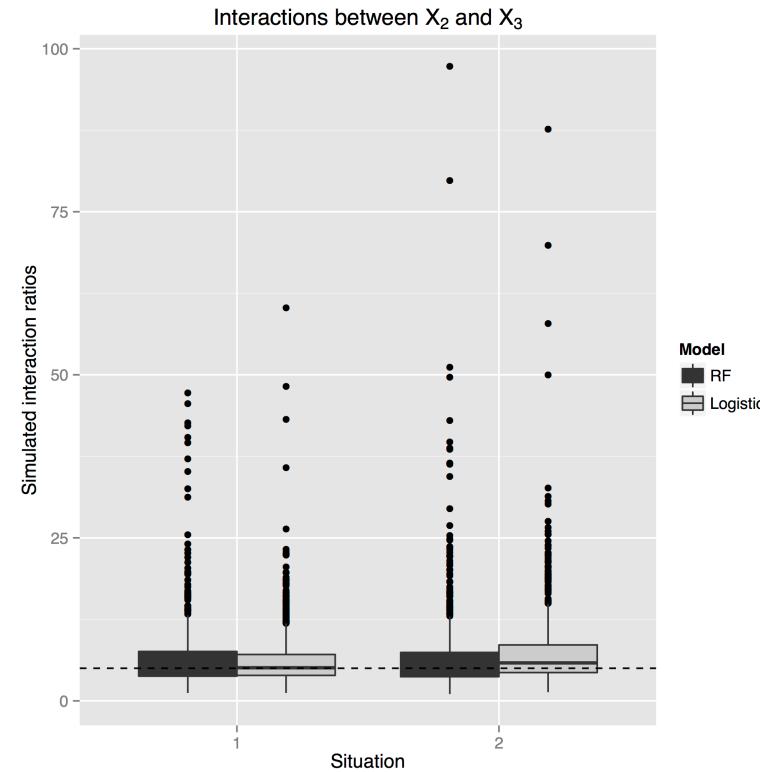
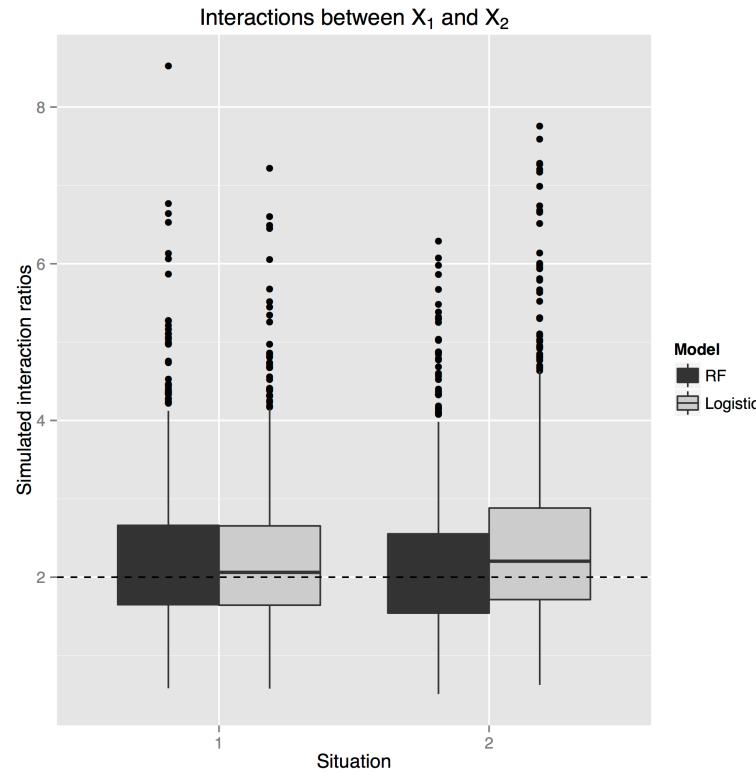
Make 4 machines to "capture landscapes" when

- $X_1 = 0, X_2 = 0 \rightarrow p_{00}$
- $X_1 = 1, X_2 = 0 \rightarrow p_{10}$
- $X_1 = 0, X_2 = 1 \rightarrow p_{01}$
- $X_1 = 1, X_2 = 1 \rightarrow p_{11}$

Now compute the appropriate contrast ( $p_{11} - p_{10} - p_{01} + p_{00}$ ) or ratio

$$\frac{p_{11}(1 - p_{10})}{p_{10}(1 - p_{11})} / \frac{p_{01}(1 - p_{00})}{p_{00}(1 - p_{01})}$$

# Risk machine



# More about interactions

The Risk Machine<sup>TM</sup> is cumbersome when you have many features.

We can actually do a faster scan of the data to find 2-way interactions

We call it the **Interactor**<sup>TM</sup>

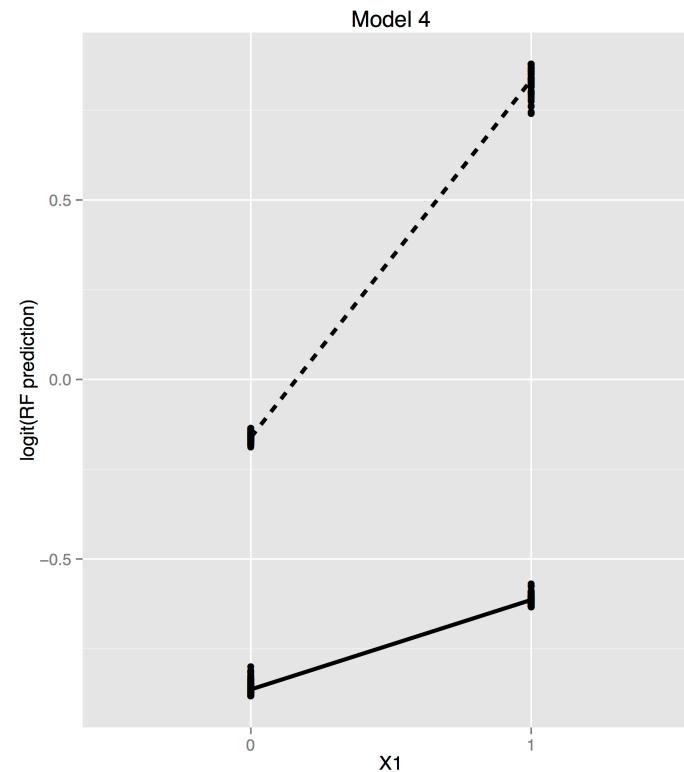
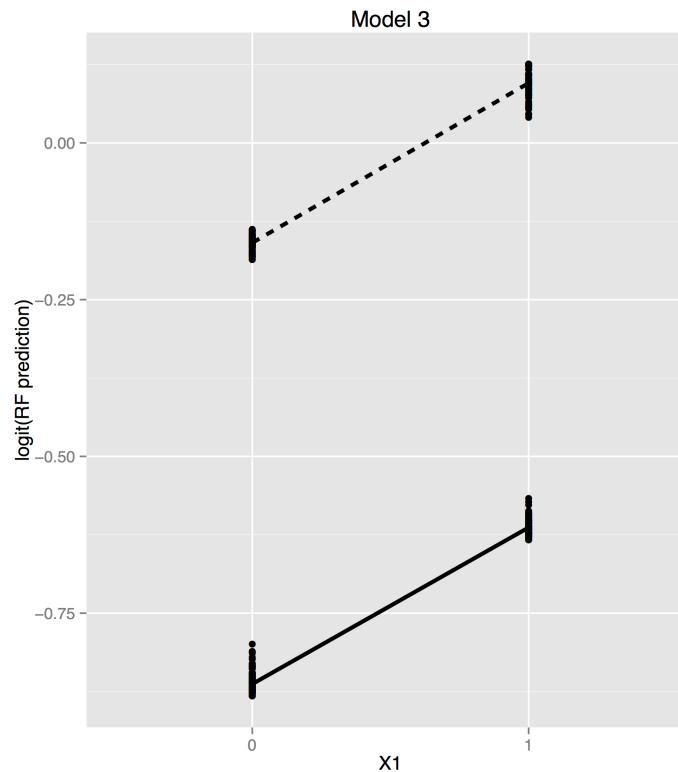
# The Interactor

We fit one PM to the data and get predicted probabilities

- Average probabilities over  $X_1=1, X_2=1 \rightarrow P_{11}$
- Average probabilities over  $X_1=0, X_2=1 \rightarrow P_{01}$
- Average probabilities over  $X_1=1, X_2=0 \rightarrow P_{10}$
- Average probabilities over  $X_1=0, X_2=0 \rightarrow P_{00}$

We can now create classical interaction plots either on natural or logit scale

# The Interactor



# The Interactor

How about when we need to scan an entire genome?

1. Run 1 PM on the dataset
2. Compute interaction contrasts for each pair of features
3. Create an interaction heatmap to find interaction hotspots
4. Drill down (MDS, more analyses)



# On to the post-parameter ecosystem